

```
In [6]: import pandas as pd
import numpy as np
import pickle
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from keras.layers import Dense, Input, LSTM, Embedding, Dropout, Activation, Conv1D, GRU, Bidirectional
from keras.layers import GlobalAveragePooling1D, GlobalMaxPooling1D, concatenate, SpatialDropout1D
from keras.models import Model, load_model
from keras.callbacks import Callback, ModelCheckpoint, EarlyStopping
from keras.optimizers import Adam
from sklearn.metrics import roc_auc_score, f1_score, precision_recall_fscore_support
import logging
import tensorflow as tf
tf.compat.v1.logging.set_verbosity(tf.compat.v1.logging.ERROR)
```

```
In [7]: # Train test split ratio of 0.3
X_train = pd.read_pickle("../Preprocessing/Data/X_train.pkl")
X_test = pd.read_pickle("../Preprocessing/Data/X_test.pkl")
y_train = pd.read_pickle("../Preprocessing/Data/y_train.pkl")
y_test = pd.read_pickle("../Preprocessing/Data/y_test.pkl")

# Url - https://dl.fbaipublicfiles.com/fasttext/vectors-english/crawl-300d-2M.vec.zip
# pre-trained FastText vector file
# word million word vectors trained on Common Crawl (600B tokens)
embedding_path = "Data/crawl-300d-2M.vec"

# no. of dimensions for each vector
embed_size = 300
# maximum number of unique words to be considered
max_features = 100000
# maximum length of a comment
max_len = 220
```

```
In [8]: # vectorize a text corpus
# turns text into sequence of space-separated sequence of words
# sequences are split into list of tokens
# they will be indexed or vectorized

# create tokenizer
token = Tokenizer(num_words=max_features, lower=True)
# fit the tokenizer on training data
# create internal vocabulary index based on word frequency
token.fit_on_texts(X_train)

# convert comment texts to their numeric counterparts
# transform each text in texts to sequence of integers
# i.e replaces each word in text with corresponding integer value from
# word_index dictionary
X_train = token.texts_to_sequences(X_train)
X_test = token.texts_to_sequences(X_test)
```

```
In [9]: # padding to make comments uniform in length
# output will be padded sequence of numbers
X_train = pad_sequences(X_train, maxlen=max_len)
X_test = pad_sequences(X_test, maxlen=max_len)
```

```
In [10]: # load FastText word embeddings
def get_coefficients(word,*arr):
    return word, np.asarray(arr, dtype='float32')

embeddings_index = dict(get_coefficients(*v.strip().split(" ")) for v in
open(embedding_path))

# create embedding matrix that contains words in our corpus and their
# corresponding values from FastText embeddings
word_index = token.word_index
nb_words = min(max_features, len(word_index))
embedding_matrix = np.zeros((nb_words, embed_size))
for word, index in word_index.items():
    if index >= max_features:
        continue
    embedding_vector = embeddings_index.get(word)
    if embedding_vector is not None:
        embedding_matrix[index] = embedding_vector
```

```
In [0]: class RocAucEvaluation(Callback):
    def __init__(self, validation_data=(), interval=1):
        super(Callback, self).__init__()

        self.interval = interval
        self.X_val, self.y_val = validation_data

    def on_epoch_end(self, epoch, logs={}):
        if epoch % self.interval == 0:
            y_pred = self.model.predict(self.X_val, verbose=0)
            score = roc_auc_score(self.y_val, y_pred)
            # results = self.model.evaluate(self.X_val, self.y_val, verbose=1)

            # precision, recall, fscore, _ = precision_recall_fscore_support(y_pred, self.y_val)
            # print("\n ROC-AUC - epoch: {:d} - score: {:.6f} - score: {:.6f} - score: {:.6f} - score: {:.6f}".format(epoch+1, score, precision, recall, fscore))
            # print("\n Results: test_loss: {:d} - test_accuracy: {:.6f}".format(results[0], results[1]))
            print("\n ROC-AUC - epoch: {:d} - score: {:.6f}".format(epoch+1, score))
```

```
In [0]: # calculating ROC_AUC score after every epoch
ra_val = RocAucEvaluation(validation_data=(X_test, y_test), interval=1)
# stop training when a monitored quantity (specified by monitor attribute) has stopped improving
# patience = number of epochs that produced the monitored quantity with no improvement after which training will be stopped
early_stop = EarlyStopping(monitor='val_loss', mode='min', patience=3)
```

```

In [0]: def build_model(lr = 0.0, lr_d = 0.0, units = 0, dr = 0.0):
    # instantiates a Keras tensor
    inp = Input(shape=(max_len,))
    # first layer
    # matrix is used to initialize weights in the Embedding layer of the model
    # trainable=False to prevent the weights from being updated during training
    x = Embedding(nb_words, embed_size, weights = [embedding_matrix], trainable = False)(inp)
    # drops entire 1D feature maps (channels) instead of individual elements
    # randomly setting a fraction rate (dr) of input units to 0 at each update, to prevent overfitting
    x1 = SpatialDropout1D(dr)(x)

    # bi-directional GRU and LSTM to keep contextual information in both directions
    x = Bidirectional(GRU(units, return_sequences = True))(x1)
    x = Conv1D(64, kernel_size = 2, padding = "valid", kernel_initializer = "he_uniform")(x)

    y = Bidirectional(LSTM(units, return_sequences = True))(x1)
    y = Conv1D(64, kernel_size = 2, padding = "valid", kernel_initializer = "he_uniform")(y)

    # to minimize overfitting
    avg_pool1 = GlobalAveragePooling1D()(x)
    max_pool1 = GlobalMaxPooling1D()(x)

    avg_pool2 = GlobalAveragePooling1D()(y)
    max_pool2 = GlobalMaxPooling1D()(y)

    # returns a single tensor by merging all the pooling layers
    x = concatenate([avg_pool1, max_pool1, avg_pool2, max_pool2])
    # Output layer which classifies a given comment into one of 6 toxic levels
    x = Dense(6, activation = "sigmoid")(x)

    # this model includes all layers required in computation of x given by inp
    model = Model(inputs = inp, outputs = x)
    # configure model for training
    model.compile(loss = "binary_crossentropy", optimizer = Adam(lr = lr, decay = lr_d), metrics = ["accuracy"])
    # trains the model for fixed number of epochs
    history = model.fit(X_train, y_train, batch_size = 128, epochs = 3, validation_data = (X_test, y_test),
                        verbose = 1, callbacks = [ra_val, early_stop])
    # model = load_model(file_path)
    return model

```

```
In [0]: # learning_rate for optimizer      = 1e-4
# learning_rate decay                    = 0
# Output dimensionality for LSTM = 128
# dropout_rate                          = 0.2
model = build_model(lr = 1e-4, lr_d = 0, units = 128, dr = 0.2)
```

Train on 39912 samples, validate on 19659 samples

Epoch 1/3

39912/39912 [=====] - 397s 10ms/step - loss:  
0.2412 - acc: 0.9144 - val\_loss: 0.1320 - val\_acc: 0.9510

ROC-AUC - epoch: 1 - score: 0.931762

Epoch 2/3

39912/39912 [=====] - 387s 10ms/step - loss:  
0.1290 - acc: 0.9505 - val\_loss: 0.1192 - val\_acc: 0.9536

ROC-AUC - epoch: 2 - score: 0.942917

Epoch 3/3

39912/39912 [=====] - 387s 10ms/step - loss:  
0.1194 - acc: 0.9531 - val\_loss: 0.1125 - val\_acc: 0.9563

ROC-AUC - epoch: 3 - score: 0.949765

```
In [0]: pickle.dump(model.history, open('Models/FastText_NN.sav', 'wb'))
```

```
In [0]: model.summary()
```

Model: "model\_1"

Layer (type) connected to	Output Shape	Param #	Connected to
input_1 (InputLayer)	(None, 220)	0	
embedding_1 (Embedding) input_1[0][0]	(None, 220, 300)	18345300	input_1
spatial_dropout1d_1 (SpatialDropout1D) embedding_1[0][0]	(None, 220, 300)	0	embedding_1
bidirectional_1 (Bidirectional) spatial_dropout1d_1[0][0]	(None, 220, 256)	329472	spatial_dropout1d_1
bidirectional_2 (Bidirectional) bidirectional_1[0][0]	(None, 220, 256)	439296	bidirectional_1
conv1d_1 (Conv1D) bidirectional_2[0][0]	(None, 219, 64)	32832	bidirectional_2
conv1d_2 (Conv1D) conv1d_1[0][0]	(None, 219, 64)	32832	conv1d_1
global_average_pooling1d_1 (GlobalAveragePooling1D) conv1d_2[0][0]	(None, 64)	0	conv1d_2
global_max_pooling1d_1 (GlobalMaxPooling1D) global_average_pooling1d_1[0][0]	(None, 64)	0	global_average_pooling1d_1
global_average_pooling1d_2 (GlobalAveragePooling1D) global_max_pooling1d_1[0][0]	(None, 64)	0	global_max_pooling1d_1
global_max_pooling1d_2 (GlobalMaxPooling1D) global_average_pooling1d_2[0][0]	(None, 64)	0	global_average_pooling1d_2
concatenate_1 (Concatenate) global_max_pooling1d_2[0][0] global_max_pooling1d_1[0][0] global_average_pooling1d_2[0][0] global_max_pooling1d_1[0][0]	(None, 256)	0	global_max_pooling1d_2 global_max_pooling1d_1 global_average_pooling1d_2 global_max_pooling1d_1

```
_max_pooling1d_2[0][0]
```

---

dense_1 (Dense)	(None, 6)	1542	concat
-----------------	-----------	------	--------

---

```
enate_1[0][0]
```

---

```
=====
Total params: 19,181,274
Trainable params: 835,974
Non-trainable params: 18,345,300
```

---

```
In [0]: results = model.evaluate(X_test, y_test, verbose=1)
```

```
print("Test Loss:", results[0])
print("Test Accuracy:", results[1])
```

```
19659/19659 [=====] - 250s 13ms/step
Test Loss: 0.11250274048393247
Test Accuracy: 0.9563049891847725
```