

```
In [7]: import pandas as pd
import numpy as np
from sklearn.metrics import accuracy_score, f1_score, roc_auc_score
import sklearn.metrics as metrics
from tqdm import tqdm
import os.path
import pickle
from keras.preprocessing.text import Tokenizer, one_hot
from keras.preprocessing.sequence import pad_sequences
from keras.models import Sequential

from keras.layers.core import Activation, Dropout, Dense
from keras.layers import Flatten, LSTM
from keras.layers import GlobalMaxPooling1D
from keras.models import Model

from sklearn.model_selection import train_test_split
from keras.utils.np_utils import to_categorical
from keras.callbacks import EarlyStopping, Callback, ModelCheckpoint
from keras.layers import Dropout, Input
from keras.layers.merge import Concatenate
from keras.layers.embeddings import Embedding

import matplotlib.pyplot as plt
import nltk
import re

from numpy import array
from numpy import asarray
from numpy import zeros

nltk.download('stopwords')
```

```
[nltk_data] Downloading package stopwords to /Users/abhay/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
```

Out[7]: True

```
In [8]: df_train = pd.read_csv('../Preprocessing/Data/train.csv')
```

```
In [9]: df_train.head()
```

```
Out[9]:
```

	id	comment_text	toxic	severe_toxic	obscene	threat	insult	identity_hate
0	0000997932d777bf	Explanation\nWhy the edits made under my usern...	0	0	0	0	0	0
1	000103f0d9c6b60f	D'aww! He matches this background colour I'm s...	0	0	0	0	0	0
2	000113f07ec002fd	Hey man, I'm really not trying to edit war. It...	0	0	0	0	0	0
3	0001b41b1c6bb37e	"\nMore\nI can't make any real suggestions on ...	0	0	0	0	0	0
4	0001d958c54c6e35	You, sir, are my hero. Any chance you remember...	0	0	0	0	0	0

```
In [10]: df_train.toxic.value_counts()
```

```
Out[10]: 0    144277
          1     15294
          Name: toxic, dtype: int64
```

```
In [11]: df_train["comment_text"][168]
```

```
Out[11]: "You should be fired, you're a moronic wimp who is too lazy to do resea
rch. It makes me sick that people like you exist in this world."
```

```
In [12]: print("Toxic:", str(df_train["toxic"][168]))
          print("Severe_toxic:", str(df_train["severe_toxic"][168]))
          print("Obscene:", str(df_train["obscene"][168]))
          print("Threat:", str(df_train["threat"][168]))
          print("Insult:", str(df_train["insult"][168]))
          print("Identity_hate:", str(df_train["identity_hate"][168]))
```

```
Toxic: 1
Severe_toxic: 0
Obscene: 0
Threat: 0
Insult: 1
Identity_hate: 0
```


```
In [13]: labels = df_train[["toxic", "severe_toxic", "obscene", "threat", "insult", "identity_hate"]]  
labels.head()
```

Out[13]:

	toxic	severe_toxic	obscene	threat	insult	identity_hate
0	0	0	0	0	0	0
1	0	0	0	0	0	0
2	0	0	0	0	0	0
3	0	0	0	0	0	0
4	0	0	0	0	0	0

```
In [14]: fig_size = plt.rcParams["figure.figsize"]  
fig_size[0] = 10  
fig_size[1] = 8  
plt.rcParams["figure.figsize"] = fig_size  
  
labels.sum(axis=0).plot.bar()
```

Out[14]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1a323d55c0>



```
In [15]: df_train.dropna(subset=['comment_text'], inplace=True)
len(df_train)
```

Out[15]: 159571

```
In [16]: df_train['comment_text'] = [entry.lower() for entry in df_train['comment_text']]
df_train.head()
```

Out[16]:

	id	comment_text	toxic	severe_toxic	obscene	threat	insult	identity_hate
0	0000997932d777bf	explanation\nwhy the edits made under my usern...	0	0	0	0	0	0
1	000103f0d9c9fb60f	d'aww! he matches this background colour i'm s...	0	0	0	0	0	0
2	000113f07ec002fd	hey man, i'm really not trying to edit war. it...	0	0	0	0	0	0
3	0001b41b1c6bb37e	"\nmore\ni can't make any real suggestions on ...	0	0	0	0	0	0
4	0001d958c54c6e35	you, sir, are my hero. any chance you remember...	0	0	0	0	0	0

```
In [17]: df_train['comment_text_lower'] = df_train['comment_text'].astype(str).str.replace('[^a-zA-Z]', ' ').str.lower()
# Remove all non-letter characters and make everything lowercase.
```

```
In [18]: stop_re = '\\b'+ '\\b|\\b'.join(nltk.corpus.stopwords.words('english'))+ '\\b'
df_train['comment_text_stop'] = df_train['comment_text_lower'].astype(str).str.replace(stop_re, '')
```

```
In [19]: df_train['comment_text_stop'].head(10)
```

Out[19]:

0	explanation	edits made	username	hardcore m...
1	aww	matches	background colour	seemingly ...
2	hey man	really	trying	edit war
3		make	real suggestions	improvement
4		sir	hero	chance
5		congratulations	well	use tools well
6			cocksucker	piss around
7	vandalism	matt shirvington	article	revert...
8	sorry	word	nonsense	offensive
9		alignment	subject	contrary

Name: comment\_text\_stop, dtype: object

```
In [20]: # tokenizing words
df_train['comment_text_final'] = df_train['comment_text_stop'].astype(str).str.split()

In [21]: X_train = pd.read_pickle('../Preprocessing/Data/X_train.pkl')
y_train = pd.read_pickle('../Preprocessing/Data/y_train.pkl')
X_test = pd.read_pickle('../Preprocessing/Data/X_test.pkl')
y_test = pd.read_pickle('../Preprocessing/Data/y_test.pkl')

In [22]: tokenizer = Tokenizer(num_words=130000)
tokenizer.fit_on_texts(X_train)

X_train = tokenizer.texts_to_sequences(X_train)
X_test = tokenizer.texts_to_sequences(X_test)

vocab_size = len(tokenizer.word_index) + 1

maxlen = 200

X_train = pad_sequences(X_train, padding='post', maxlen=maxlen)
X_test = pad_sequences(X_test, padding='post', maxlen=maxlen)

In [23]: vocab_size

Out[23]: 61152

In [25]: embeddings_dictionary = dict()

glove_file = open('Data/glove.6B.100d.txt', encoding="utf8")

for line in glove_file:
    records = line.split()
    word = records[0]
    vector_dimensions = asarray(records[1:], dtype='float32')
    embeddings_dictionary[word] = vector_dimensions
glove_file.close()

embedding_matrix = zeros((vocab_size, 100))
for word, index in tokenizer.word_index.items():
    embedding_vector = embeddings_dictionary.get(word)
    if embedding_vector is not None:
        embedding_matrix[index] = embedding_vector

In [26]: deep_inputs = Input(shape=(maxlen,))
embedding_layer = Embedding(vocab_size, 100, weights=[embedding_matrix],
trainable=False)(deep_inputs)
LSTM_Layer_1 = LSTM(128)(embedding_layer)
dense_layer_1 = Dense(6, activation='sigmoid')(LSTM_Layer_1)
model = Model(inputs=deep_inputs, outputs=dense_layer_1)

model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['acc'])
```

```
In [27]: class RocAucEvaluation(Callback):
    def __init__(self, validation_data=(), interval=1):
        super(Callback, self).__init__()

        self.interval = interval
        self.X_val, self.y_val = validation_data

    def on_epoch_end(self, epoch, logs={}):
        if epoch % self.interval == 0:
            y_pred = self.model.predict(self.X_val, verbose=0)
            score = roc_auc_score(self.y_val, y_pred)
            print("\n ROC-AUC - epoch: {:d} - score: {:.6f}".format(epoch+1, score))
ra_val = RocAucEvaluation(validation_data=(X_test, y_test), interval=1)
```

```
In [28]: ra_val = RocAucEvaluation(validation_data=(X_test, y_test), interval=1)
early_stop = EarlyStopping(monitor='val_loss', mode='min', patience=5)
```

```
In [29]: print(model.summary())
```

Model: "model\_1"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	(None, 200)	0
-----		
embedding_1 (Embedding)	(None, 200, 100)	6115200
-----		
lstm_1 (LSTM)	(None, 128)	117248
-----		
dense_1 (Dense)	(None, 6)	774
=====		
Total params: 6,233,222		
Trainable params: 118,022		
Non-trainable params: 6,115,200		
-----		
None		

```
In [31]: history = None
if os.path.isfile('Models/lstm_glove.sav'):
    history = pickle.load(open('Models/lstm_glove.sav', 'rb'))
else:
    history = model.fit(X_train, y_train, batch_size = 128, epochs = 5,
        validation_data = (X_test, y_test),
        verbose = 1, callbacks = [ra_val, early_stop], v
        alidation_split = 0.2)
    pickle.dump(history, open('Models/lstm_glove.sav', 'wb'))
```

Train on 39912 samples, validate on 19659 samples

Epoch 1/5

39912/39912 [=====] - 139s 3ms/step - loss: 0.2891 - acc: 0.9016 - val\_loss: 0.2733 - val\_acc: 0.9026

ROC-AUC - epoch: 1 - score: 0.509909

Epoch 2/5

39912/39912 [=====] - 138s 3ms/step - loss: 0.2725 - acc: 0.9024 - val\_loss: 0.2542 - val\_acc: 0.9029

ROC-AUC - epoch: 2 - score: 0.859188

Epoch 3/5

39912/39912 [=====] - 132s 3ms/step - loss: 0.1915 - acc: 0.9267 - val\_loss: 0.1918 - val\_acc: 0.9206

ROC-AUC - epoch: 3 - score: 0.880922

Epoch 4/5

39912/39912 [=====] - 135s 3ms/step - loss: 0.1585 - acc: 0.9393 - val\_loss: 0.1392 - val\_acc: 0.9470

ROC-AUC - epoch: 4 - score: 0.922844

Epoch 5/5

39912/39912 [=====] - 138s 3ms/step - loss: 0.1382 - acc: 0.9472 - val\_loss: 0.1323 - val\_acc: 0.9498

ROC-AUC - epoch: 5 - score: 0.927736

```
In [32]: score = model.evaluate(X_test, y_test, verbose=1)
```

```
print("Test Loss:", score[0])
```

```
print("Test Accuracy:", score[1])
```

19659/19659 [=====] - 21s 1ms/step

Test Loss: 0.13227544613025263

Test Accuracy: 0.949776828289032


```
In [33]: plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])


plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])

plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```







In [0]: